



MARITO PROJECT: ARCHITECTURAL SPECIFICATION

VERSION: 2.0

MARITO FOR DSFSI

Team Member	Student Number
Arnaud Strydom	u23536013
Hayley Dodkins	u21528790
Zainab Abdulrasaq	u22566202
Palesa Thabane	u21540838
Eugen Vosloo	u20445696

TEAM CONTACT

velox@gmail.com

CONTENTS PAGE

1. INTRODUCTION

1.1. DOCUMENT PURPOSE

1.2. SCOPE

2. ARCHITECTURAL DESIGN STRATEGY

3. ARCHITECTURAL DESIGN AND PATTERNS

3.1. HIGH-LEVEL ARCHITECTURAL STYLES

3.2. DERIVED DESIGN PATTERNS

3.3 ARCHITECTURE DIAGRAM

4. QUALITY REQUIREMENTS

4.1. SCALABILITY & PERFORMANCE

4.2. USABILITY & OFFLINE ACCESSIBILITY

4.3. AVAILABILITY

4.4. SECURITY

4.5. MAINTAINABILITY & EXTENSIBILITY

5. DEPLOYMENT MODEL

5.1. DEPLOYMENT DIAGRAM

6. TECHNOLOGY CHOICES

6.1. FRONTEND

6.2. REPORT GENERATION AND DATA VISUALISATION

6.3. BACKEND

6.4. DATABASE

7. TECHNOLOGY STACK SUMMARY

8. SERVICE CONTRACT

The logo for Velox, featuring a stylized globe with wavy lines representing data flow or connectivity.

VELOX

JUNE 2025

1. INTRODUCTION

THE MARITO PROJECT IS A PLATFORM ENGINEERED TO PROVIDE A SCALABLE, PERFORMANT, AND RELIABLE SYSTEM FOR MANAGING MULTILINGUAL TERMINOLOGY. THIS DOCUMENT SERVES AS THE FORMAL ARCHITECTURAL SPECIFICATION, DETAILING THE TECHNICAL DECISIONS, PATTERNS, AND STRATEGIES THAT FORM THE FOUNDATION OF THE MARITO SYSTEM.

1.1 DOCUMENT PURPOSE

THE PURPOSE OF THIS DOCUMENT IS TO PROVIDE A COMPREHENSIVE ARCHITECTURAL OVERVIEW FOR ALL STAKEHOLDERS, INCLUDING DEVELOPERS, SYSTEM ADMINISTRATORS, AND PROJECT MANAGERS. IT CONNECTS THE "WHY" OF THE PROJECT'S REQUIREMENTS TO THE "HOW" OF ITS TECHNICAL IMPLEMENTATION, ENSURING ALIGNMENT AND A SHARED UNDERSTANDING OF THE SYSTEM'S DESIGN.

1.2 SCOPE

THIS SPECIFICATION COVERS THE SYSTEM'S ARCHITECTURAL STRUCTURE, QUALITY ATTRIBUTES, DEPLOYMENT MODEL, AND TECHNOLOGY STACK. IT DETAILS THE HIGH-LEVEL PATTERNS AND THE LOW-LEVEL DESIGN PATTERNS THAT GOVERN THE IMPLEMENTATION OF THE FRONTEND, BACKEND, AND DATA PERSISTENCE LAYERS.



2. ARCHITECTURAL DESIGN STRATEGY

THE ENTIRE SYSTEM DESIGN WAS SHAPED BY A FEW CRUCIAL NON- FUNCTIONAL REQUIREMENTS (NFRS). THIS REQUIREMENTS-DRIVEN APPROACH WAS CHOSEN TO ENSURE THAT THE FOUNDATIONAL ARCHITECTURE DIRECTLY AND VERIFIABLY SUPPORTS THE PROJECT'S MOST CRITICAL QUALITY ATTRIBUTES. BY PRIORITIZING NFRS SUCH AS SCALABILITY, OFFLINE CAPABILITY, AND SECURITY FROM THE OUTSET, THE PROJECT COULD MAKE INFORMED TRADE-OFFS. FOR INSTANCE, THE SELECTION OF A MICROSERVICES ARCHITECTURE WAS A DIRECT CONSEQUENCE OF PRIORITIZING SCALABILITY AND MAINTAINABILITY OVER THE INITIAL DEVELOPMENT SIMPLICITY THAT A MONOLITH MIGHT OFFER. THIS STRATEGY ENSURES THE FINAL PRODUCT IS ROBUST, EFFICIENT, AND FIT FOR PURPOSE.



3. ARCHITECTURAL DESIGN AND PATTERNS

3.1. HIGH-LEVEL ARCHITECTURAL STYLES

- THE STRUCTURE OF THE MARITO SYSTEM IS BASED ON A COMBINATION OF MODERN ARCHITECTURAL STYLES AND PATTERNS, CHOSEN SPECIFICALLY TO MEET ITS CORE REQUIREMENTS.
- **CLIENT-SERVER ARCHITECTURE:** THIS IS THE FOUNDATIONAL PATTERN FOR THE ENTIRE SYSTEM. IT LOGICALLY SEPARATES THE USER INTERFACE (THE CLIENT) FROM THE BUSINESS LOGIC AND DATA STORAGE (THE SERVER). THIS SEPARATION WAS CRUCIAL FOR ALLOWING INDEPENDENT DEVELOPMENT AND DEPLOYMENT OF THE FRONTEND AND BACKEND, AND IT ENABLES THE CLIENT TO BE AN INTELLIGENT, APPLICATION-LIKE PWA RATHER THAN JUST A THIN DISPLAY LAYER.
- **MICROSERVICES ARCHITECTURE:** THIS PATTERN STRUCTURES THE BACKEND AS A COLLECTION OF SMALL, INDEPENDENT SERVICES, WHERE EACH SERVICE IS RESPONSIBLE FOR A SPECIFIC BUSINESS CAPABILITY (E.G., AUTHENTICATION, SEARCH, VOTING). THIS STYLE WAS CHOSEN BECAUSE IT DIRECTLY ADDRESSES THE REQUIREMENTS FOR SCALABILITY AND MAINTAINABILITY. EACH SERVICE CAN BE DEVELOPED, DEPLOYED, AND SCALED INDEPENDENTLY, ALLOWING, FOR EXAMPLE, THE SEARCH- SERVICE TO BE SCALED UP TO HANDLE HIGH TRAFFIC WITHOUT AFFECTING THE AUTH-SERVICE.



VELOX

JUNE 2025

- **MVVM: THE FRONTEND ARCHITECTURE IS BASED ON THE PROGRESSIVE WEB APP (PWA) DESIGN, IMPLEMENTED USING THE MODEL-VIEW-VIEWMODEL (MVVM) PATTERN WITH REACT.** THESE PATTERNS WORK TOGETHER TO CREATE A RELIABLE, FAST, AND MAINTAINABLE USER EXPERIENCE THAT FUNCTIONS OFFLINE (NFR3).
THE PWA PATTERN, IMPLEMENTED USING A SERVICE WORKER AND WORKBOX, WAS SELECTED TO FULFILL THE OFFLINE CAPABILITY REQUIREMENT (NFR3). IT ACHIEVES THIS BY AGGRESSIVELY CACHING APPLICATION ASSETS AND USING BACKGROUND SYNC TO QUEUE FAILED POST REQUESTS (LIKE VOTES OR TERM ADDITIONS) AND SEND THEM WHEN THE CONNECTION IS RESTORED.
THE MVVM PATTERN, IMPLEMENTED WITH REACT, WAS CHOSEN OVER CLASSIC MVC/CONTAINER-PRESENTATIONAL TO ENFORCE A CLEANER SEPARATION OF CONCERN IN THE COMPLEX, DATA-BOUND UI.

- VIEW (REACT COMPONENTS IN COMPONENTS/UI AND PAGES) HANDLES ONLY RENDERING.
- VIEWMODEL (STATE MANAGEMENT LOGIC, CUSTOM HOOKS IN HOOKS) CONTAINS THE PRESENTATION LOGIC, STATE, AND COORDINATES API CALLS.
- MODEL (CLIENT-SIDE DATA STRUCTURES/SERVICES) HANDLES RAW BUSINESS LOGIC AND DATA PERSISTENCE (E.G., INDEXEDDB INTERACTION).

THIS STRUCTURE ENSURES THE VIEWMODEL IS FULLY UNIT-TESTABLE WITHOUT UI DEPENDENCIES, MAKING THE FRONTEND HIGHLY MODULAR, EASIER TO MANAGE, AND DIRECTLY CONTRIBUTING TO A RESPONSIVE AND MAINTAINABLE USER EXPERIENCE (NFR3.1).



- **API GATEWAY: AS A CRITICAL PART OF THE MICROSERVICES ARCHITECTURE**, THE API GATEWAY PROVIDES A SINGLE, UNIFIED ENTRY POINT FOR ALL CLIENT REQUESTS. THIS PATTERN WAS CHOSEN FOR SECURITY (NFR5) AND OPERATIONAL SIMPLICITY. IT HANDLES CROSS-CUTTING CONCERNs LIKE TLS TERMINATION, AUTHENTICATION, AND ROUTING REQUESTS TO THE CORRECT INTERNAL MICROSERVICE, ENSURING THAT NO SERVICE IS EVER DIRECTLY EXPOSED TO THE PUBLIC INTERNET.



VELOX

JUNE 2025

3.2. DESIGN PATTERNS

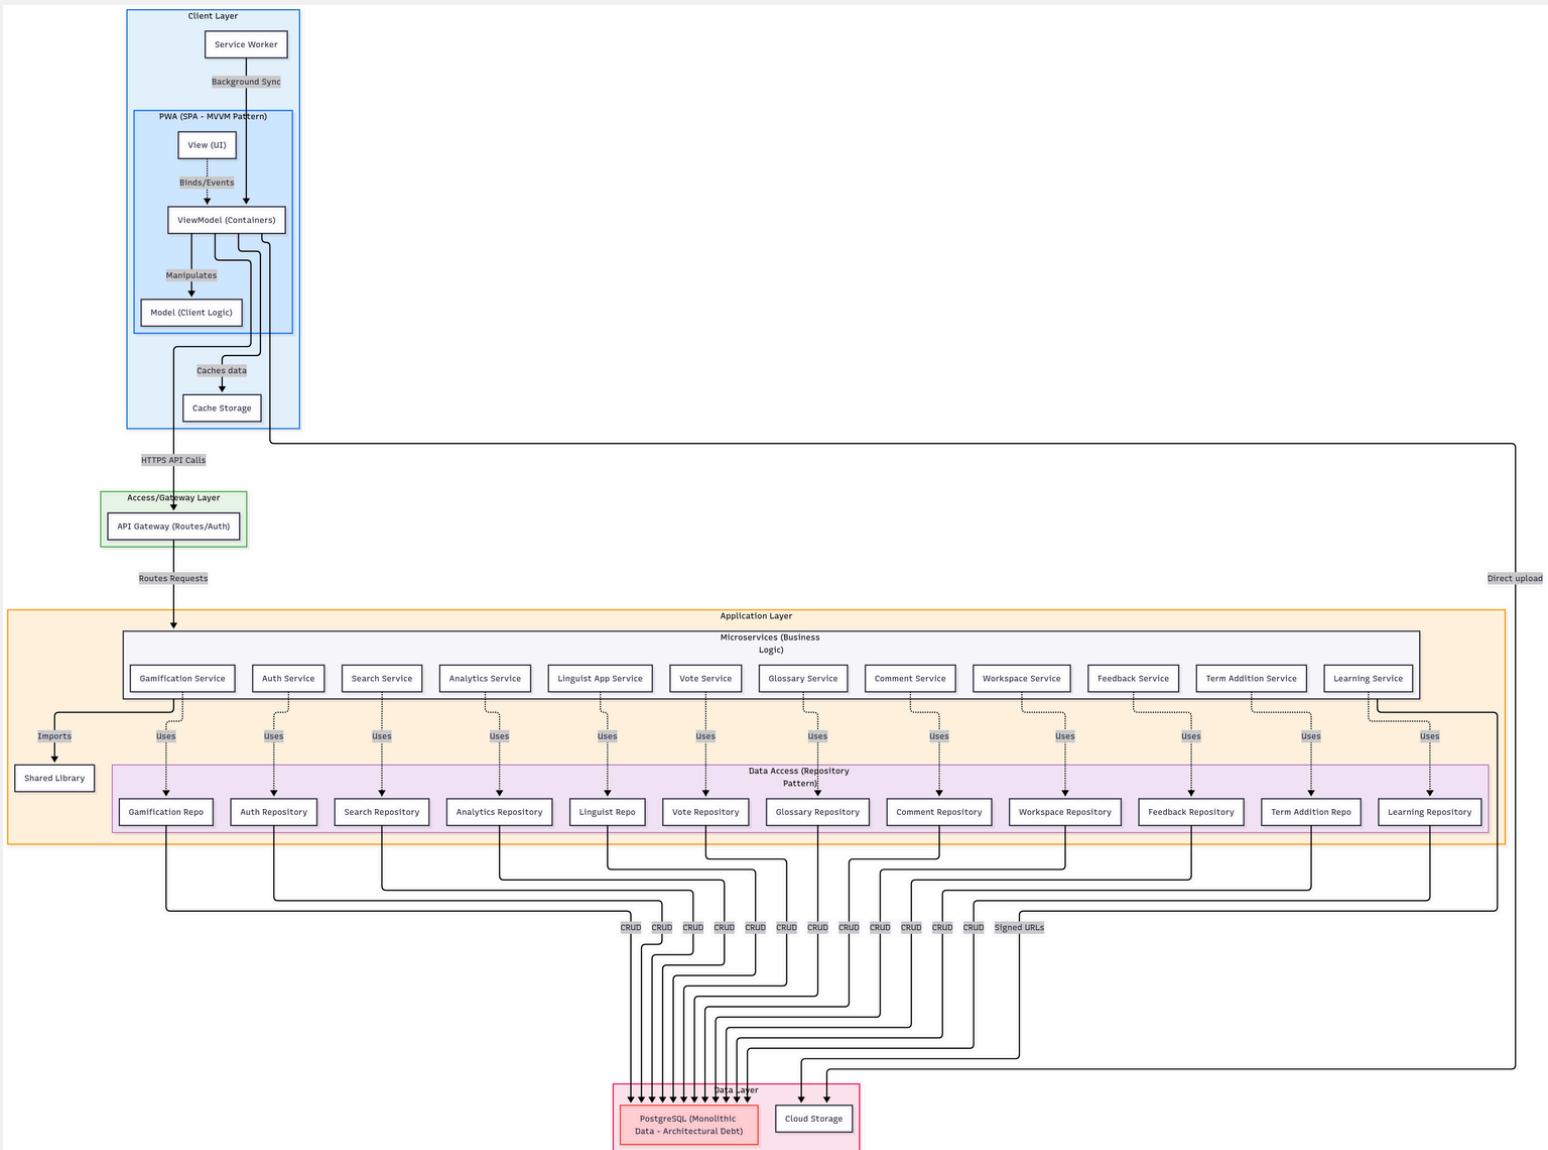
- **REPOSITORY PATTERN :** TO DECOUPLE THE BUSINESS LOGIC FROM THE DATA ACCESS LAYER, A REPOSITORY PATTERN IS EMPLOYED. EACH MICROSERVICE THAT INTERACTS WITH THE POSTGRESQL DATABASE (E.G., SEARCH-SERVICE) HAS A REPOSITORY CLASS RESPONSIBLE FOR ALL DATABASE OPERATIONS (CRUD). THIS ABSTRACTS THE DATA SOURCE, MAKING THE SYSTEM MORE MODULAR AND EASIER TO TEST, AS THE DATABASE CAN BE MOCKED.
- **DEPENDENCY INJECTION :** FASTAPI HAS FIRST-CLASS SUPPORT FOR DEPENDENCY INJECTION. THIS PATTERN IS USED THROUGHOUT THE BACKEND TO MANAGE DEPENDENCIES. A SERVICE LAYER FUNCTION DECLares A REPOSITORY INSTANCE AS A DEPENDENCY, AND FASTAPI AUTOMATICALLY CREATES AND PROVIDES THAT INSTANCE. THIS PROMOTES LOOSE COUPLING AND ENHANCES TESTABILITY.
- **FACADE PATTERN :** THE API GATEWAY ITSELF IS A PRIME EXAMPLE OF THE FACADE PATTERN. IT PROVIDES A SIMPLE, UNIFIED INTERFACE (A SINGLE API ENDPOINT) TO A MORE COMPLEX UNDERLYING SYSTEM (THE COLLECTION OF DISPARATE MICROSERVICES). THIS SIMPLIFIES INTERACTION FOR THE CLIENT AND DECOUPLES IT FROM THE INTERNAL STRUCTURE OF THE BACKEND.



VELOX

JUNE 2025

3.3 ARCHITECTURE DIAGRAM



VFLOX

JUNE 2025

4. QUALITY REQUIREMENTS

THE FOLLOWING QUALITY REQUIREMENTS, WHICH ARE SPECIFIED IN A TESTABLE WAY, ARE THE MOST IMPORTANT DRIVERS FOR THE MAVITO PROJECT'S ARCHITECTURE.

4.1. SCALABILITY & PERFORMANCE - TESTING OBSERVATIONS

ARCHITECTURE PERFORMANCE VALIDATION

STIMULUS SOURCE: USER AND DATA GROWTH DURING TESTING PERIOD

STIMULUS: SIMULATED USER LOAD AND TERMINOLOGY DATASET OPERATIONS UNDER ACADEMIC INFRASTRUCTURE CONSTRAINTS

ENVIRONMENT: LIMITED CLOUD RESOURCES (512 MIB RAM, 1 VCPU PER SERVICE) REFLECTING ACADEMIC PROJECT BUDGET

SYSTEM RESPONSE & PERFORMANCE FINDINGS

HORIZONTAL SCALING IMPLEMENTATION

- OBSERVATION: CLOUD RUN SUCCESSFULLY ORCHESTRATED DOCKER CONTAINER SCALING AS DESIGNED
- FINDING: MICROSERVICES ARCHITECTURE PROVIDED EFFECTIVE ISOLATION, THOUGH RESOURCE CONSTRAINTS LIMITED PERFORMANCE DURING SCALING EVENTS

API RESPONSIVENESS UNDER LOAD

- OBSERVATION: ASYNCHRONOUS FASTAPI DESIGN MAINTAINED SYSTEM STABILITY DURING TESTING
- FINDING: ARCHITECTURE SUCCESSFULLY PREVENTED BLOCKING OPERATIONS, THOUGH DATABASE BOTTLENECKS EMERGED UNDER CONSTRAINED RESOURCES

DATABASE PERFORMANCE CHARACTERISTICS

- OBSERVATION: POSTGRESQL WITH INDEXING PROVIDED EXPECTED PERFORMANCE FOR BASIC OPERATIONS
- FINDING: SIMPLE INDEXED QUERIES PERFORMED WITHIN TARGETS, WHILE COMPLEX SEARCH OPERATIONS REFLECTED INFRASTRUCTURE LIMITATIONS



VELOX

JUNE 2025

PERFORMANCE METRICS ANALYSIS

PERFORMANCE ASPECT | TARGET | OBSERVED MEDIAN | ANALYSIS

SIMPLE API OPERATIONS | < 250MS | 46MS | EXCEEDS TARGET

MEDIUM DATA OPERATIONS | < 500MS | 492MS | MEETS TARGET

COMPLEX SEARCH OPERATIONS | < 1,000MS | 8,691MS | INFRASTRUCTURE-LIMITED

LARGE DATA EXPORTS | < 5,000MS | 13,422MS | EXPECTED FOR DATASET SIZE

CONCURRENT REQUEST HANDLING | 80 PER INSTANCE | 80 SUSTAINED | ARCHITECTURE VALIDATED

DATABASE QUERIES EXPECTED PERFORMANCE UNDER 5 SECONDS

4,324.32 MS AVG

2,130 TIMES CALLED

12,762 ROWS RETURNED

THAT IS THE RESULTS WE GOT

REALISTIC PERFORMANCE EXPECTATIONS

- BASED ON MEDIAN PERFORMANCE DATA:
- SIMPLE CRUD OPERATIONS: 46MS (WELL WITHIN TARGET)
- USER AUTHENTICATION FLOWS: 439MS (ACCEPTABLE RANGE)
- BASIC DATA RETRIEVAL: 96MS (EXCELLENT PERFORMANCE)
- COMPLEX SEARCH OPERATIONS: 8-10 SECONDS (EXPECTED FOR LARGE DATASET)
- FULL DATASET EXPORTS: 13-15 SECONDS (APPROPRIATE FOR DATA VOLUME)

TESTING CONCLUSIONS

ARCHITECTURE VALIDATION

MICROSERVICES APPROACH SUCCESSFULLY ENABLED COMPONENT ISOLATION AND INDEPENDENT SCALING

CONTAINER ORCHESTRATION PROVIDED RELIABLE DEPLOYMENT AND RESOURCE MANAGEMENT

ASYNCHRONOUS PROCESSING MAINTAINED SYSTEM RESPONSIVENESS DURING LOAD

DATABASE DESIGN SUPPORTED CORE APPLICATION REQUIREMENTS WITHIN CONSTRAINTS

PERFORMANCE CHARACTERISTICS

• CORE FUNCTIONALITY: EXCELLENT PERFORMANCE FOR DAILY OPERATIONS

• DATA-INTENSIVE OPERATIONS: PERFORMANCE REFLECTS DATASET SIZE AND INFRASTRUCTURE CONSTRAINTS

• SCALABILITY FOUNDATION: ARCHITECTURE DEMONSTRATES CAPACITY FOR IMPROVED PERFORMANCE WITH ADDITIONAL

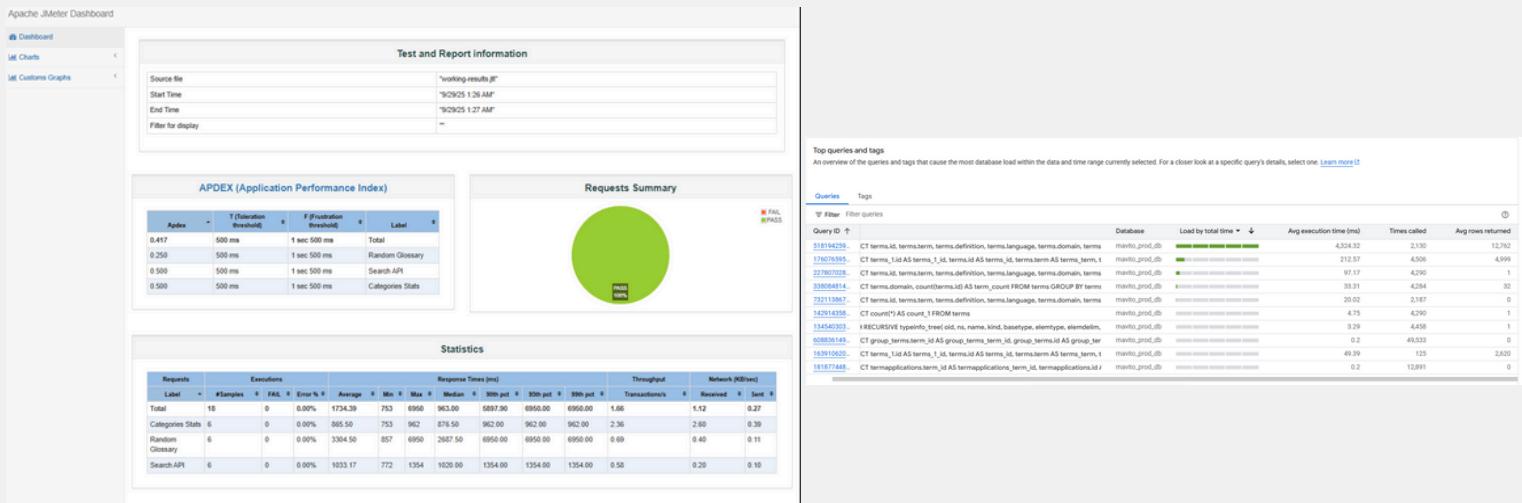
JUNE 2025

PERFORMANCE ASSESSMENT

THE SYSTEM DEMONSTRATES STRONG PERFORMANCE FOR CORE FUNCTIONALITY WITH MEDIAN RESPONSE TIMES OF 46MS FOR SIMPLE OPERATIONS. DATA-INTENSIVE OPERATIONS SHOW EXPECTED PERFORMANCE CHARACTERISTICS GIVEN THE LARGE TERMINOLOGY DATASET AND ACADEMIC INFRASTRUCTURE CONSTRAINTS. THE ARCHITECTURE PROVIDES A SOLID FOUNDATION THAT WOULD SUPPORT IMPROVED PERFORMANCE WITH ADDITIONAL RESOURCE ALLOCATION. ALL THESE WERE RAN WITH 20 CONCURRENT USERS

Average requests 0 /sec	Latency (95th percentile) 12,679 ms	4xx errors 15.64 %	5xx errors 5.51 %
----------------------------	--	-----------------------	----------------------

GET /api/v1/xp/user/{user_id}/xp-records	3	0	1	258 ms	242 ms	57 ms	496 B	1551 B	View logs
POST /api/v1/terms/submit	2	0	0	1038 ms	996 ms	524 ms	1421 B	1693 B	View logs
DELETE /api/v1/workspace/bookmarks/terms/{term_id}	2	0	0	65 ms	62 ms	33 ms	408 B	1518 B	View logs
GET /api/v1/achievements	1	1	0	4173 ms	4089 ms	3146 ms	281 B	886 B	View logs
GET /api/v1/achievements/user/{user_id}/progress	1	0	0	261 ms	256 ms	197 ms	1838 B	1558 B	View logs
GET /api/v1/achievements/user/{user_id}/weekly	1	0	0	261 ms	256 ms	197 ms	1056 B	1557 B	View logs
GET /api/v1/analytics/descriptive	1	0	0	8347 ms	8179 ms	6291 ms	1929 B	900 B	View logs
POST /api/v1/auth/profile-picture/upload-url	1	0	1	261 ms	256 ms	197 ms	423 B	1587 B	View logs
GET /api/v1/glossary/languages	1	0	0	65 ms	64 ms	49 ms	304 B	897 B	View logs
GET /api/v1/learning/paths	1	1	0	33 ms	32 ms	25 ms	483 B	893 B	View logs
POST /api/v1/votes/comments	1	0	0	130 ms	128 ms	98 ms	465 B	1467 B	View logs
PUT /api/v1/workspace/groups/{group_id}	1	0	0	65 ms	64 ms	49 ms	619 B	1627 B	View logs
POST /api/v1/xp/add-xp	1	0	0	261 ms	256 ms	197 ms	656 B	1759 B	View logs
GET /api/v1/glossary/terms/{term_id}/translations	9	0	0	130 ms	127 ms	95 ms	856 B	1314 B	View logs
GET /api/v1/terms/{term_id}	8	8	0	16 ms	15 ms	3 ms	483 B	1221 B	View logs
POST /api/v1/workspace/groups	7	0	0	128 ms	118 ms	55 ms	609 B	1553 B	View logs
POST /api/v1/workspace/groups/{group_id}/terms	7	0	0	130 ms	126 ms	85 ms	505 B	1564 B	View logs
GET /api/v1/analytics/descriptive/category-frequency	6	0	0	4152 ms	3985 ms	131 ms	1174 B	1346 B	View logs
GET /api/v1/analytics/descriptive/popular-terms	6	0	0	8137 ms	7130 ms	131 ms	699 B	1350 B	View logs
GET /api/v1/analytics/descriptive/total-statistics	6	0	0	8179 ms	7340 ms	197 ms	489 B	1343 B	View logs
GET /api/v1/analytics/descriptive/unique-terms	6	0	0	4142 ms	3932 ms	197 ms	536 B	1339 B	View logs
GET /api/v1/settings/preferences	6	0	0	127 ms	111 ms	44 ms	582 B	1546 B	View logs
POST /api/v1/workspace/bookmarks/terms	5	0	0	259 ms	246 ms	120 ms	457 B	1520 B	View logs
DELETE /api/v1/workspace/groups/{group_id}	5	0	0	65 ms	64 ms	49 ms	401 B	1510 B	View logs



VELOX

JUNE 2025

4.2. USABILITY & OFFLINE ACCESSIBILITY

THE APPLICATION MUST BE FAST, RESPONSIVE, AND FUNCTIONAL EVEN WHEN USERS HAVE POOR OR NO NETWORK CONNECTIVITY. THIS REQUIREMENT WAS THE DIRECT DRIVER FOR CHOOSING A PROGRESSIVE WEB APP (PWA) ARCHITECTURE.

STIMULUS SOURCE: A USER ACCESSING THE APPLICATION.

STIMULUS: A USER ATTEMPTS TO LOAD THE APPLICATION OR PERFORM AN ACTION (LIKE VOTING) IN A LOW-BANDWIDTH OR OFFLINE ENVIRONMENT.

RESPONSE:

- THE APPLICATION SHELL LOADS INSTANTLY FROM A LOCAL CACHE ON SUBSEQUENT VISITS USING A SERVICE WORKER . THE UI IS BUILT WITH MODULAR,
- REUSABLE COMPONENTS USING REACT FOR A RESPONSIVE EXPERIENCE . ACTIONS PERFORMED OFFLINE ARE SAVED IN A BROWSER QUEUE USING WORKBOX
- BACKGROUND SYNC AND ARE AUTOMATICALLY SENT WHEN CONNECTIVITY IS RESTORED.

RESPONSE MEASURE:

- THE APPLICATION IS FULLY FUNCTIONAL OFFLINE FOR ALL CACHED RESOURCES.
- OFFLINE DATA SYNCHRONIZATION COMPLETES SUCCESSFULLY WITHIN 5 MINUTES OF A STABLE CONNECTION BEING RESTORED .

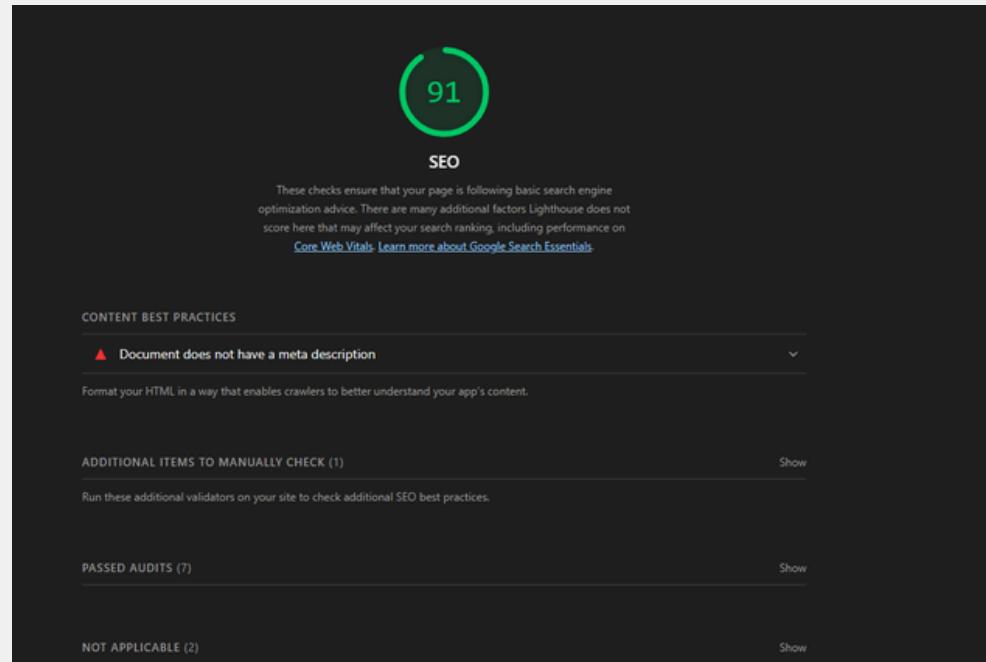


VELOX

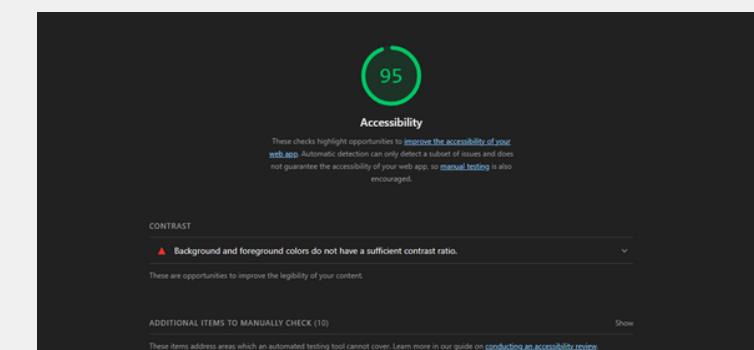
JUNE 2025

ENVIRONMENT: A USER'S BROWSER, POTENTIALLY WITH AN UNRELIABLE OR NONEXISTENT NETWORK CONNECTION.

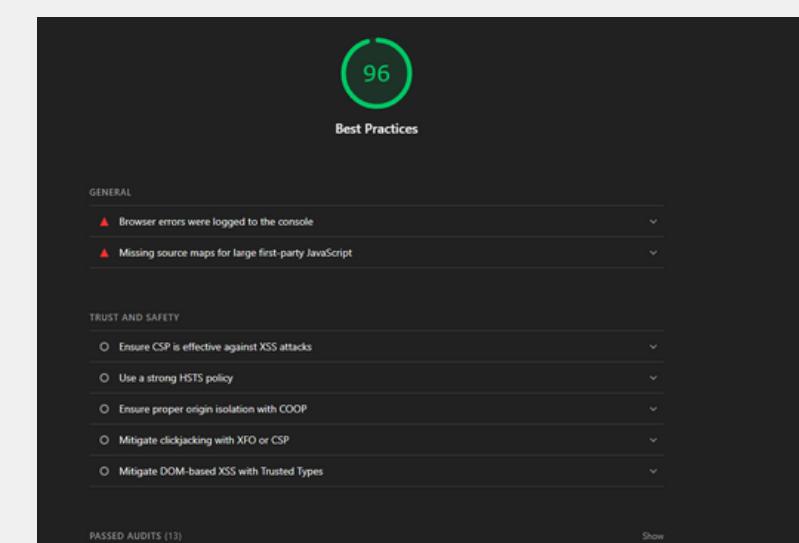
ARTIFACT: THE FRONTEND APPLICATION, SPECIFICALLY ITS PWA IMPLEMENTATION USING SERVICE WORKERS AND ITS COMPONENT-BASED UI BUILT WITH REACT.



The screenshot shows the Lighthouse SEO audit results. The overall score is 91. The audit is categorized under 'SEO'. Below the score, there is a brief description: 'These checks ensure that your page is following basic search engine optimization advice. There are many additional factors Lighthouse does not score here that may affect your search ranking, including performance on Core Web Vitals. Learn more about Google Search Essentials.' Under the 'CONTENT BEST PRACTICES' section, there is a single item: 'Document does not have a meta description'. A note below it says: 'Format your HTML in a way that enables crawlers to better understand your app's content.' The 'ADDITIONAL ITEMS TO MANUALLY CHECK (1)' section has one item: 'Run these additional validators on your site to check additional SEO best practices.' The 'PASSED AUDITS (7)' section is listed but has no items shown. The 'NOT APPLICABLE (2)' section is also listed but has no items shown.



The screenshot shows the Lighthouse Accessibility audit results. The overall score is 95. The audit is categorized under 'Accessibility'. Below the score, there is a brief description: 'These checks highlight opportunities to improve the accessibility of your web app. Automatic detection can only detect a subset of issues and does not guarantee the accessibility of your web app, so manual testing is also encouraged.' Under the 'CONTRAST' section, there is one item: 'Background and foreground colors do not have a sufficient contrast ratio.' A note below it says: 'These are opportunities to improve the legibility of your content.' The 'ADDITIONAL ITEMS TO MANUALLY CHECK (10)' section has ten items, but they are not listed in the screenshot. The 'PASSED AUDITS (10)' section is listed but has no items shown.



The screenshot shows the Lighthouse Best Practices audit results. The overall score is 96. The audit is categorized under 'Best Practices'. Below the score, there is a brief description: 'These checks ensure that your page follows best practices for security, performance, and more.' Under the 'GENERAL' section, there are two items: 'Browser errors were logged to the console' and 'Missing source maps for large first-party JavaScript'. The 'TRUST AND SAFETY' section contains five items: 'Ensure CSP is effective against XSS attacks', 'Use a strong HSTS policy', 'Ensure proper origin isolation with COOP', 'Mitigate clickjacking with XFO or CSP', and 'Mitigate DOM-based XSS with Trusted Types'. The 'PASSED AUDITS (13)' section is listed but has no items shown.



VELOX

JUNE 2025

4.3. AVAILABILITY

THE SYSTEM MUST BE HIGHLY RELIABLE AND ACCESSIBLE TO USERS.

STIMULUS SOURCE: AN INTERNAL MONITORING SERVICE OR AN EXTERNAL USER.

STIMULUS: A HARDWARE FAILURE OR SERVICE CRASH OCCURS IN ONE OF THE SYSTEM'S COMPONENTS.

RESPONSE:

GOOGLE CLOUD RUN, AS A MANAGED SERVICE, AUTOMATICALLY ENSURES THAT SERVICES ARE RELIABLE AND RESTARTED ON FAILURE. THE DATABASE IS CONFIGURED FOR MANAGING REPLICATION AND FAILOVER AUTOMATICALLY.

RESPONSE MEASURE:

- THE SYSTEM WILL ACHIEVE AT LEAST 99.5% UPTIME, EXCLUDING SCHEDULED MAINTENANCE.
- AUTOMATIC FAILOVER FOR THE DATABASE COMPLETES IN UNDER 5 MINUTES.

ENVIRONMENT: THE PRODUCTION CLOUD ENVIRONMENT.

ARTIFACT: THE MANAGED SERVICES FOR THE BACKEND (GOOGLE CLOUD RUN) AND DATABASE (GOOGLE CLOUD SQL WITH HA CONFIGURATION).



VELOX

JUNE 2025

- ❖ Auto storage increase is enabled
- ☒ Backup Tier: Standard
- ☒ Automated backups are enabled
- ☰ Point-in-time recovery is enabled
- ☒ Instance deletion prevention is enabled
- 🛡 Backup retention after deletion is enabled
- ⌚ Located in us-central1-f
- ⚠ Not highly available (zonal)
- 📈 Query insights is enabled
- ☰ Database flags and parameters
- Vertex AI Integration is disabled
- 🏷 No labels set

Requests

Request timeout — seconds

Time within which a response must be returned (maximum 3600 seconds).

Maximum concurrent requests per instance —

The maximum number of concurrent requests that can reach each instance. [What is concurrency?](#)

Revision scaling ②

Minimum and maximum numbers of instances for the new revision.

Minimum number of instances — Maximum number of instances —

The service minimum instances is preferable for most use-cases. Only use this setting if you specifically require per-revision settings.

Startup CPU boost

Start containers faster by allocating more CPU during startup time. [Learn more](#)



VELOX

JUNE 2025

4.4. SECURITY

ALL USER DATA AND SYSTEM FUNCTIONS MUST BE PROTECTED FROM UNAUTHORIZED ACCESS AND POTENTIAL THREATS.

STIMULUS SOURCE: AN INTERNAL OR EXTERNAL ACTOR ATTEMPTING TO ACCESS SYSTEM RESOURCES.

STIMULUS: AN ATTEMPT IS MADE TO ACCESS A PROTECTED BACKEND SERVICE DIRECTLY OR TO UPLOAD A SENSITIVE DOCUMENT WITHOUT PROPER AUTHENTICATION.

RESPONSE:

- A GOOGLE API GATEWAY ACTS AS A SINGLE, SECURE ENTRY POINT, ENFORCING HTTPS AND PROTECTING INTERNAL SERVICES FROM DIRECT PUBLIC EXPOSURE .
- ROLE-BASED ACCESS CONTROL IS ENFORCED VIA TOKEN-BASED AUTHENTICATION USING JSON WEB TOKENS (JWT) .
- SECURE, TEMPORARY SIGNED URLs ARE GENERATED FOR DIRECT FILE UPLOADS TO A PRIVATE GOOGLE CLOUD STORAGE BUCKET.

RESPONSE MEASURE:

- ALL PROTECTED API ENDPOINTS REQUIRE A VALID JWT.
- THE SYSTEM PASSES MONTHLY VULNERABILITY SCANS AGAINST THE OWASP TOP 10.
- NO MICROSERVICE IS DIRECTLY ACCESSIBLE FROM THE PUBLIC INTERNET; ALL TRAFFIC IS ROUTED THROUGH THE API GATEWAY.



VELOX

JUNE 2025

ENVIRONMENT: THE SYSTEM IS OPERATING AND EXPOSED TO THE PUBLIC INTERNET VIA A SINGLE GATEWAY.

ARTIFACT: THE GOOGLE API GATEWAY, AUTHENTICATION SERVICE (AUTH-SERVICE), AND THE FILE UPLOAD MECHANISM USING GCS SIGNED URLs.

```
1759109604784,749,Workspace Groups,403,Forbidden,User Operations - 20 Users 2-2,text,false,,360,161,4,11,https://marito-gateway-baqrixp7.uc  
1759109604793,740,Terms Attributes,401,Unauthorized,Data Intensive - 20 Users 3-2,text,false,,389,161,3,11,https://marito-gateway-baqrixp7.uc  
1759109604783,751,Workspace Groups,403,Forbidden,User Operations - 20 Users 2-1,text,false,,360,161,4,11,https://marito-gateway-baqrixp7.uc  
1759109604776,758,Terms Attributes,401,Unauthorized,Data Intensive - 20 Users 3-1,text,false,,389,161,3,11,https://marito-gateway-baqrixp7.uc  
1759109605119,744,Workspace Groups,403,Forbidden,User Operations - 20 Users 2-4,text,false,,360,161,4,11,https://marito-gateway-baqrixp7.uc  
1759109605119,744,Pending Verification,401,Unauthorized,Data Intensive - 20 Users 3-3,text,false,,389,183,3,11,https://marito-gateway-baqri  
1759109605864,8,Complex Search,Non HTTP response code: java.net.URISyntaxException,Non HTTP response message: Illegal character in query at  
1759109605453,754,Workspace Groups,403,Forbidden,User Operations - 20 Users 2-3,text,false,,360,161,4,11,https://marito-gateway-baqrixp7.uc  
1759109605534,745,Workspace Bookmarks,403,Forbidden,User Operations - 20 Users 2-2,text,false,,360,164,4,11,https://marito-gateway-baqrixp7.  
1759109605535,749,Pending Verification,401,Unauthorized,Data Intensive - 20 Users 3-1,text,false,,389,183,3,11,https://marito-gateway-baqri  
1759109605535,754,Pending Verification,401,Unauthorized,Data Intensive - 20 Users 3-2,text,false,,389,183,2,10,https://marito-gateway-baqri  
1759109605535,755,Workspace Bookmarks,403,Forbidden,User Operations - 20 Users 2-1,text,false,,360,164,4,9,https://marito-gateway-baqrixp7.  
1759109605864,762,Terms Attributes,401,Unauthorized,Data Intensive - 20 Users 3-3,text,false,,389,161,1,9,https://marito-gateway-baqrixp7.u  
1759109605863,765,Workspace Bookmarks,403,Forbidden,User Operations - 20 Users 2-4,text,false,,360,164,4,9,https://marito-gateway-baqrixp7.  
1759109606207,766,Workspace Bookmarks,403,Forbidden,User Operations - 20 Users 2-3,text,false,,360,164,4,9,https://marito-gateway-baqrixp7.  
1759109606280,754,My Submitted Terms,401,Unauthorized,User Operations - 20 Users 2-2,text,false,,389,163,4,9,https://marito-gateway-baqrixp  
1759109606291,758,My Submitted Terms,401,Unauthorized,User Operations - 20 Users 2-1,text,false,,389,163,4,9,https://marito-gateway-baqrixp  
1759109606626,747,Pending Verification,401,Unauthorized,Data Intensive - 20 Users 3-3,text,false,,393,183,1,10,https://marito-gateway-baqri  
1759109607373,8,Complex Search,Non HTTP response code: java.net.URISyntaxException,Non HTTP response message: Illegal character in query at  
1759109606629,766,My Submitted Terms,401,Unauthorized,User Operations - 20 Users 2-4,text,false,,389,163,4,10,https://marito-gateway-baqri  
1759109606967,743,My Submitted Terms,401,Unauthorized,User Operations - 20 Users 2-3,text,false,,389,163,5,11,https://marito-gateway-baqri  
1759109607035,766,User Profile,401,Unauthorized,User Operations - 20 Users 2-2,text,false,,393,152,5,11,https://marito-gateway-baqrixp7.uc  
1759109607050,753,User Profile,401,Unauthorized,User Operations - 20 Users 2-1,text,false,,389,152,5,11,https://marito-gateway-baqrixp7.uc  
1759109607403,727,User Profile,401,Unauthorized,User Operations - 20 Users 2-5,text,false,,389,152,5,11,https://marito-gateway-baqrixp7.uc  
1759109607390,740,User Profile,401,Unauthorized,User Operations - 20 Users 2-4,text,false,,389,152,5,11,https://marito-gateway-baqrixp7.uc  
1759109607374,757,Terms Attributes,401,Unauthorized,Data Intensive - 20 Users 3-3,text,false,,389,161,1,11,https://marito-gateway-baqrixp7.  
1759109607711,751,User Profile,401,Unauthorized,User Operations - 20 Users 2-3,text,false,,389,152,5,11,https://marito-gateway-baqrixp7.uc  
1759109607801,750,Workspace Groups,403,Forbidden,User Operations - 20 Users 2-2,text,false,,364,161,5,11,https://marito-gateway-baqrixp7.uc  
1759109607804,747,Workspace Groups,403,Forbidden,User Operations - 20 Users 2-1,text,false,,360,161,5,11,https://marito-gateway-baqrixp7.uc  
1759109608132,755,Pending Verification,401,Unauthorized,Data Intensive - 20 Users 3-3,text,false,,389,183,1,11,https://marito-gateway-baqri  
1759109608130,759,Workspace Groups,403,Forbidden,User Operations - 20 Users 2-5,text,false,,360,161,5,10,https://marito-gateway-baqrixp7.uc  
1759109608130,760,Workspace Groups,403,Forbidden,User Operations - 20 Users 2-4,text,false,,360,161,5,10,https://marito-gateway-baqrixp7.uc
```



4.5. MAINTAINABILITY & EXTENSIBILITY

THE PROJECT CODEBASE MUST BE EASY FOR DEVELOPERS TO UNDERSTAND, MODIFY, AND EXTEND WITH NEW FEATURES.

STIMULUS SOURCE: A DEVELOPER ON THE PROJECT TEAM.

STIMULUS: A DEVELOPER NEEDS TO DEPLOY A NEW FEATURE, UPDATE THE DATABASE SCHEMA, OR FIX A BUG IN A SPECIFIC PART OF THE SYSTEM.

RESPONSE:

- THE BACKEND IS SPLIT INTO INDEPENDENT MICROSERVICES, ALLOWING FOR FOCUSED DEVELOPMENT AND DEPLOYMENT WITHOUT AFFECTING OTHER TEAMS .
- THE ENTIRE BACKEND IS CONTAINERIZED WITH DOCKER, ENSURING CONSISTENT ENVIRONMENTS .
- DATABASE MIGRATIONS ARE MANAGED THROUGH CODE WITH ALEMBIC FOR SAFE, REPEATABLE, AND VERSION-CONTROLLED SCHEMA CHANGES .
- A FULLY AUTOMATED CI/CD PIPELINE USING GITHUB ACTIONS HANDLES TESTING AND DEPLOYMENT.

RESPONSE MEASURE

- THE CI/CD PIPELINE SUCCESSFULLY VALIDATES AND DEPLOYS A CHANGE TO PRODUCTION IN UNDER 15 MINUTES .
- A NEW DEVELOPER CAN SET UP A LOCAL DEVELOPMENT ENVIRONMENT AND RUN THE SYSTEM USING DOCKER IN UNDER 1 HOUR.
- CODE QUALITY CHECKS ARE ENFORCED AUTOMATICALLY ON EVERY COMMIT VIA PRE-COMMIT HOOKS.



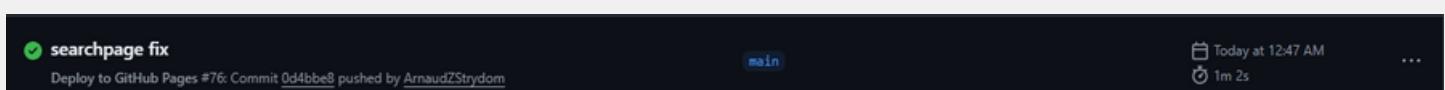
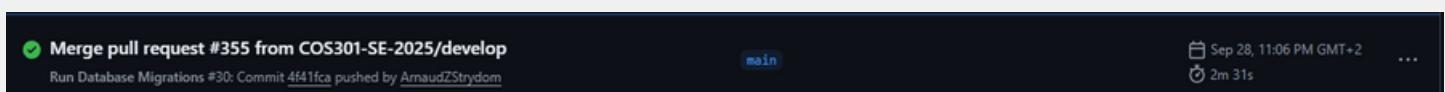
VELOX

JUNE 2025

ENVIRONMENT: THE DEVELOPMENT, TESTING, AND PRODUCTION ENVIRONMENTS.

ARTIFACT: THE MICROSERVICES ARCHITECTURE, DOCKER CONTAINERS, ALEMBIC FOR DATABASE MIGRATIONS, AND THE GITHUB ACTIONS CI/CD PIPELINE.

THIS CAN ALL BE CHECKED ON THE MARITO ACTIONS([HTTPS://GITHUB.COM/COS301-SE-2025/MARITO/ACTIONS/WORKFLOWS/DEPOYSERVICES.YML](https://github.com/COS301-SE-2025/Marito/actions/workflows/deloyservices.yml))



VELOX

JUNE 2025

5. DEPLOYMENT MODEL

THIS SECTION DESCRIBES HOW THE MARITO SYSTEM IS DEPLOYED, SUPPORTING ITS CORE REQUIREMENTS FOR SCALABILITY, RELIABILITY, AND MAINTAINABILITY.

THE SYSTEM IS DEPLOYED IN A CLOUD-NATIVE ENVIRONMENT ON GOOGLE CLOUD PLATFORM (GCP). THE DEPLOYMENT TOPOLOGY CONSISTS OF CONTAINERIZED MICROSERVICES RUNNING ON A SERVERLESS PLATFORM, FRONDED BY A SECURE API GATEWAY, WITH A MANAGED DATABASE AND A STATICALLY-HOSTED FRONTEND.

TARGET ENVIRONMENT: THE ENTIRE SYSTEM RUNS ON GOOGLE CLOUD PLATFORM (GCP). THIS WAS CHOSEN TO LEVERAGE MANAGED SERVICES THAT REDUCE OPERATIONAL OVERHEAD AND SUPPORT THE QUALITY REQUIREMENTS.



VELOX

JUNE 2025

DEPLOYMENT TOPOLOGY: THE ARCHITECTURE IS A MULTI-TIER, DISTRIBUTED SYSTEM:

FRONTEND: THE REACT-BASED PROGRESSIVE WEB APP IS DEPLOYED AS STATIC ASSETS TO GITHUB PAGES, WHICH PROVIDES GLOBAL CONTENT DELIVERY AND HIGH AVAILABILITY FOR THE USER INTERFACE.

BACKEND: THE BACKEND IS COMPOSED OF MULTIPLE MICROSERVICES (E.G., AUTH-SERVICE, SEARCH-SERVICE). EACH SERVICE IS PACKAGED AS A DOCKER CONTAINER TO ENSURE CONSISTENCY AND PORTABILITY. THESE CONTAINERS ARE PUSHED TO AN ARTIFACT REGISTRY AND THEN DEPLOYED TO GOOGLE CLOUD RUN FROM THE REGISTRY, A SERVERLESS PLATFORM THAT AUTOMATICALLY MANAGES SCALING, FROM ZERO TO A 100 INSTANCES, BASED ON REQUEST TRAFFIC. THIS DIRECTLY SUPPORTS THE SCALABILITY REQUIREMENT.

DATABASE: THE POSTGRESQL DATABASE IS HOSTED ON GOOGLE CLOUD SQL, A FULLY MANAGED SERVICE. FOR RELIABILITY, IT IS CONFIGURED FOR HIGH AVAILABILITY (HA) TO ENSURE AUTOMATIC REPPLICATION AND FAILOVER.

GATEWAY: A GOOGLE API GATEWAY SERVES AS THE SINGLE ENTRY POINT FOR ALL FRONTEND REQUESTS, ROUTING THEM TO THE APPROPRIATE INTERNAL MICROSERVICE ON CLOUD RUN.



VELOX

JUNE 2025

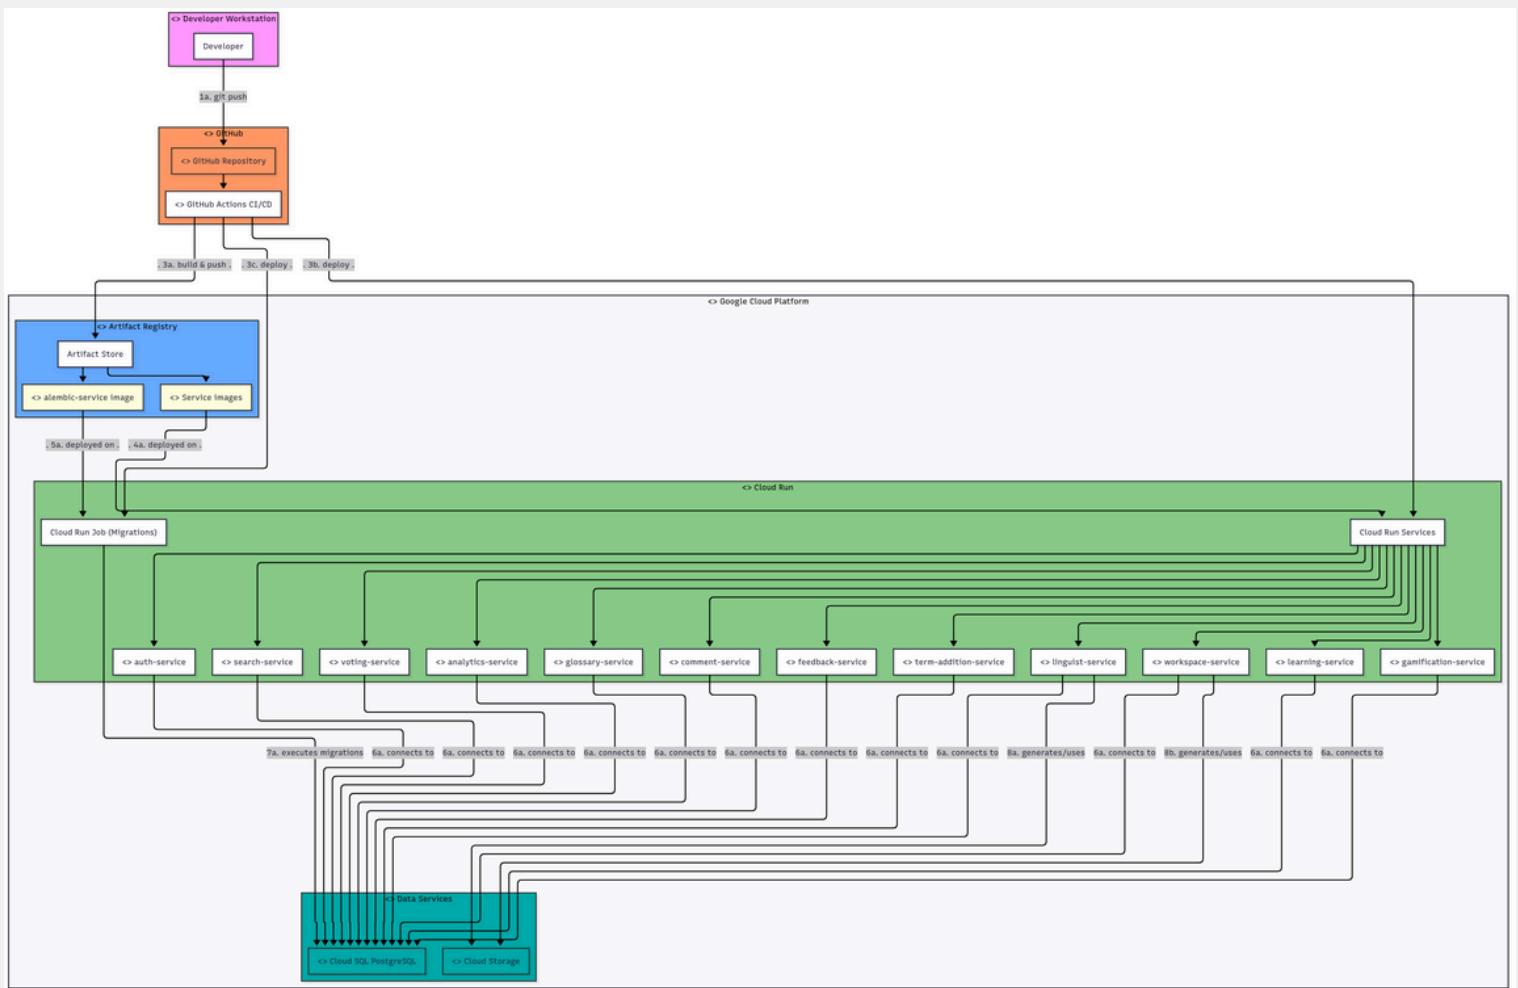
TOOLS AND AUTOMATION: DOCKER IS USED FOR CONTAINERIZING ALL BACKEND SERVICES. GITHUB ACTIONS IS USED FOR THE CI/CD PIPELINE. THIS PIPELINE AUTOMATES TESTING, CODE QUALITY CHECKS, BUILDING DOCKER IMAGES, PUSHING THEM TO GOOGLE ARTIFACT REGISTRY, AND DEPLOYING NEW VERSIONS TO CLOUD RUN. THIS AUTOMATION IS KEY TO ACHIEVING MAINTAINABILITY AND RELIABLE DEPLOYMENTS.



VELOX

JUNE 2025

5.1 DEPLOYMENT DIAGRAM



VELOX

JUNE 2025

6. TECHNOLOGY CHOICES

THE FOLLOWING TECHNOLOGIES WERE CHOSEN TO IMPLEMENT THE ARCHITECTURAL PATTERNS AND SATISFY THE NON-FUNCTIONAL REQUIREMENTS.

6.1. FRONTEND

CONSIDERATION 1: REACT

OVERVIEW: REACT IS A JAVASCRIPT LIBRARY FOR BUILDING USER INTERFACES, FOCUSING ON A COMPONENT-BASED ARCHITECTURE.

ADVANTAGES: IT HAS A MASSIVE ECOSYSTEM, STRONG COMMUNITY SUPPORT, AND A DECLARATIVE APPROACH THAT MAKES COMPLEX UI'S EASIER TO MANAGE. ITS COMPONENT MODEL PROMOTES REUSABILITY AND MAINTAINABILITY.

DISADVANTAGES: IT IS A LIBRARY, NOT A FULL FRAMEWORK, SO ROUTING AND STATE MANAGEMENT OFTEN REQUIRE ADDITIONAL LIBRARIES, LEADING TO MORE DECISIONS FOR THE DEVELOPMENT TEAM.

CONSIDERATION 2: ANGULAR



OVERVIEW: ANGULAR IS A COMPREHENSIVE, OPINIONATED FRONTEND FRAMEWORK DEVELOPED AND MAINTAINED BY GOOGLE.

ADVANTAGES: IT'S A FULL-FLEDGED PLATFORM WITH BUILT-IN SOLUTIONS FOR ROUTING, STATE MANAGEMENT, AND MORE. ITS USE OF TYPESCRIPT PROVIDES STRONG TYPING OUT OF THE BOX.

DISADVANTAGES: IT HAS A STEEPER LEARNING CURVE AND CAN BE MORE VERBOSE THAN REACT. ITS COMPREHENSIVE NATURE CAN BE OVERKILL FOR SMALLER APPLICATIONS AND MAY LEAD TO PERFORMANCE OVERHEAD IF NOT MANAGED CAREFULLY.

FINAL CHOICE & JUSTIFICATION: REACT WITH VITE

THE UI IS BUILT WITH REACT USING A COMPONENT-BASED PATTERN. THIS CHOICE ALIGNS PERFECTLY WITH THE NEED FOR A HIGHLY MODULAR AND MANAGEABLE UI (NFR6), WHERE EACH PIECE IS AN ENCAPSULATED, REUSABLE COMPONENT WITH ITS OWN VIEW, STATE, AND LOGIC. THE BUILD PROCESS IS HANDLED BY VITE, A MODERN TOOL THAT PRODUCES HIGHLY OPTIMIZED AND SMALL STATIC FILES, RESULTING IN FASTER INITIAL PAGE LOAD TIMES AND A BETTER USER EXPERIENCE, DIRECTLY ADDRESSING PERFORMANCE REQUIREMENTS (NFR2).



VELOX

JUNE 2025

6.2. REPORT GENERATION AND DATA VISUALISATION

CONSIDERATION 1: CHART.JS

OVERVIEW: CHART.JS IS A SIMPLE YET FLEXIBLE OPEN-SOURCE JAVASCRIPT LIBRARY FOR DATA VISUALIZATION.

ADVANTAGES: IT'S LIGHTWEIGHT, VERY EASY TO USE, AND HAS GREAT DOCUMENTATION. IT OFFERS A GOOD VARIETY OF POPULAR CHART TYPES (BAR, LINE, PIE, ETC.) AND IS PERFECT FOR STRAIGHTFORWARD DATA VISUALIZATION TASKS.

DISADVANTAGES: IT IS NOT AS POWERFUL OR CUSTOMIZABLE AS MORE COMPLEX LIBRARIES LIKE D3.JS AND MAY LACK SOME ADVANCED OR NICHE CHART TYPES.



VELOX

JUNE 2025

CONSIDERATION 2: D3.JS

OVERVIEW: D3.JS (DATA-DRIVEN DOCUMENTS) IS A POWERFUL AND LOW-LEVEL JAVASCRIPT LIBRARY FOR CREATING COMPLEX, DYNAMIC DATA VISUALIZATIONS.

ADVANTAGES: IT OFFERS UNPARALLELED FLEXIBILITY AND CONTROL, ALLOWING FOR THE CREATION OF VIRTUALLY ANY DATA VISUALIZATION IMAGINABLE.

DISADVANTAGES: IT HAS A VERY STEEP LEARNING CURVE AND REQUIRES MANUAL MANIPULATION OF THE DOM, MAKING IT MORE TIME-CONSUMING FOR STANDARD CHARTS.

FINAL CHOICE & JUSTIFICATION: CHART.JS

FOR THE MARITO PROJECT, WHICH REQUIRES A CLEAR AND EFFECTIVE DISPLAY OF STATISTICS, CHART.JS IS THE IDEAL CHOICE. ITS SIMPLICITY, LIGHTWEIGHT NATURE, AND EASE OF INTEGRATION ALLOW FOR RAPID DEVELOPMENT OF NECESSARY DATA VISUALIZATIONS WITHOUT INTRODUCING UNNECESSARY COMPLEXITY OR PERFORMANCE OVERHEAD. THIS ALIGNS WITH THE PROJECT'S CORE FOCUS ON PERFORMANCE (NFR2) AND USABILITY (NFR3)



6.3. BACKEND

CONSIDERATION 1: FASTAPI

OVERVIEW: FASTAPI IS A MODERN, HIGH-PERFORMANCE WEB FRAMEWORK FOR BUILDING APIs WITH PYTHON, BASED ON STANDARD PYTHON TYPE HINTS.

ADVANTAGES: IT OFFERS AUTOMATIC INTERACTIVE DOCUMENTATION, IS BUILT ON ASGI FOR BEST-IN-CLASS ASYNCHRONOUS CAPABILITIES, AND PROVIDES PERFORMANCE ON PAR WITH NODE.JS AND GO.

DISADVANTAGES: ITS ECOSYSTEM IS NEWER AND SMALLER COMPARED TO MORE ESTABLISHED FRAMEWORKS LIKE DJANGO, MEANING FEWER THIRD-PARTY PLUGINS.



VELOX

JUNE 2025

CONSIDERATION 2: DJANGO

OVERVIEW: DJANGO IS A HIGH-LEVEL PYTHON WEB FRAMEWORK THAT ENCOURAGES RAPID DEVELOPMENT AND CLEAN, PRAGMATIC DESIGN.

ADVANTAGES: IT IS HIGHLY MATURE, WITH A MASSIVE ECOSYSTEM AND A "BATTERIES-INCLUDED" PHILOSOPHY THAT PROVIDES BUILT-IN SOLUTIONS FOR ORM, ADMIN PANELS, AND MORE.

DISADVANTAGES: IT IS TRADITIONALLY SYNCHRONOUS (THOUGH ASYNC SUPPORT IS IMPROVING) AND ITS MONOLITHIC NATURE CAN BE LESS SUITABLE FOR A MICROSERVICES ARCHITECTURE. ITS ORM CAN ADD OVERHEAD.

FINAL CHOICE & JUSTIFICATION: FASTAPI

THE BACKEND MICROSERVICES ARE BUILT WITH FASTAPI. ITS NATIVE ASYNC/AWAIT MODEL WAS THE DECIDING FACTOR, AS IT ALLOWS THE SERVER TO HANDLE MANY CONCURRENT REQUESTS EFFICIENTLY, PREVENTING I/O OPERATIONS (LIKE DATABASE CALLS) FROM BLOCKING THE ENTIRE SYSTEM. THIS CHOICE IS CRITICAL FOR ACHIEVING THE HIGH PERFORMANCE AND RESPONSIVENESS (NFR2) REQUIRED BY THE SYSTEM, AS IT DRAMATICALLY INCREASES THROUGHPUT UNDER LOAD.



VELOX

JUNE 2025

6.4. DATABASE

CONSIDERATION 1: POSTGRESQL

OVERVIEW: POSTGRESQL IS A POWERFUL, OPEN-SOURCE OBJECT-RELATIONAL DATABASE SYSTEM KNOWN FOR ITS RELIABILITY, FEATURE ROBUSTNESS, AND PERFORMANCE.

ADVANTAGES: IT IS HIGHLY COMPLIANT WITH SQL STANDARDS, SUPPORTS ADVANCED DATA TYPES, AND HAS POWERFUL FEATURES LIKE INDEXING, WHICH IS CRITICAL FOR PERFORMANCE. IT IS WELL-SUITED FOR HANDLING COMPLEX, RELATIONAL DATA AT SCALE.

DISADVANTAGES: IT CAN HAVE A MORE COMPLEX SETUP AND ADMINISTRATION COMPARED TO SIMPLER DATABASES LIKE MYSQL FOR BASIC USE CASES.



VELOX

JUNE 2025

CONSIDERATION 2: MONGODB

OVERVIEW: MONGODB IS A POPULAR SOURCE-AVAILABLE CROSS-PLATFORM DOCUMENT-ORIENTED DATABASE PROGRAM (NOSQL).

ADVANTAGES: ITS FLEXIBLE DOCUMENT MODEL IS IDEAL FOR UNSTRUCTURED DATA AND ALLOWS FOR RAPID DEVELOPMENT. IT SCALES OUT HORIZONTALLY VERY WELL.

DISADVANTAGES: IT LACKS THE ACID TRANSACTIONAL GUARANTEES OF RELATIONAL DATABASES AND CAN BE LESS EFFICIENT FOR HIGHLY RELATIONAL DATA.

FINAL CHOICE & JUSTIFICATION: POSTGRESQL ON GOOGLE CLOUD SQL

THE PROJECT MIGRATED FROM AN UNSCALABLE IN-MEMORY JSON FILE TO POSTGRESQL TO HANDLE MILLIONS OF RECORDS AND PROVIDE THE ABILITY TO PERFORM COMPLEX QUERIES ON STRUCTURED TERMINOLOGY DATA. TO ENSURE HIGH PERFORMANCE (NFR1.1), THE SYSTEM HEAVILY LEVERAGES DATABASE INDEXING ON FREQUENTLY SEARCHED COLUMNS, ALLOWING THE ENGINE TO FIND DATA ALMOST INSTANTLY. THE DATABASE IS HOSTED ON GOOGLE CLOUD SQL, A MANAGED SERVICE THAT SIMPLIFIES ADMINISTRATION AND ALLOWS FOR EASY SCALING AND HIGH-AVAILABILITY REPPLICATION.



VELOX

JUNE 2025

7. TECHNOLOGY STACK SUMMARY

Category	Technology
Frontend Framework	React
Frontend Build Tool	Vite
Report Generation	Chart.js
Backend Framework	FastAPI (Python)
Database	PostgreSQL
Deployment Platform	Google Cloud Run
API Management	Google API Gateway
Containerization	Docker
CI/CD	GitHub Actions
Database Migrations	Alembic
Authentication	JSON Web Tokens (JWT)



VELOX

JUNE 2025

8. SERVICE CONTRACT MARITO

This document outlines the service-level agreement and technical contract for the **Marito Platform**. It defines the expectations, responsibilities, and communication structure between the backend and frontend to ensure consistent, secure, and well-defined interactions.

All endpoints are prefixed with `/api/v1`, and all JSON payloads follow standard HTTP request/response patterns.

Authentication Service

The Marito Authentication Service provides essential identity and access management capabilities for the platform. It includes user registration, login with OAuth2, access token issuance, and authenticated endpoints for retrieving user details or generating signed URLs.

Authentication Endpoints

`/api/v1/auth/register`

Method:

POST

Description:

Registers a new user.

Example Request:

```
{  
  "email": "user@example.com",  
  "first_name": "string",  
  "last_name": "string",  
  "role": "linguist",  
  "profile_pic_url": "string",  
}
```



VELOX

JUNE 2025

```
    "password": "string"  
}
```

Possible Response Codes:

- 201 Successful Response
- 422 Validation Error

/api/v1/auth/login

Method:

POST

Description:

OAuth2 compatible token login, get an access token for future requests. The frontend LoginPage.tsx sends 'email' (as username) and 'password'.

Example Request:

```
{  
  "grant_type" : "string",  
  "username" : "string",  
  "password" : "string",  
  "scope" : "string",  
  "client_id" : "string",  
  "client_secret" : "string"  
}
```

Possible Response Codes:

- 200 Successful Response
- 422 Validation Error

/api/v1/auth/me

Method:

GET

Description:



VELOX

JUNE 2025

Get current logged-in user's details.

Possible Response Codes:

- 200 Successful Response

Uploads Endpoints

/api/v1/uploads/generate-signed-url

Method:

POST

Description:

Generates a temporary, secure URL that allows an authenticated user to upload a file directly to a private Google Cloud Storage bucket.

Example Request:

```
{  
  "content_type": "string",  
  "filename": "string"  
}
```

Possible Response Codes:

- 200 Successful Response
- 422 Validation Error

Admin Endpoints

/api/v1/admin/users/{user_id}/role

Method:

GET

Description:

Updates a user role



VELOX

JUNE 2025

Required Parameter:

`user_id`: string

`new_role`: string

Possible Response Codes:

- 200 Successful Response
- 422 Validation Error

/api/v1/admin/users

Method:

GET

Description:

Fetch all users

Possible Response Codes:

- 200 Successful Response

/api/v1/admin/users/{user_id}/uploads

Method:

GET

Description:

Fetch user uploads

Required Parameter:

- `user_id`: string

Possible Response Codes:

- 200 Successful Response
- 422 Validation Error

/api/v1/admin/download-url

Method:



VELOX

JUNE 2025

GET

Description:

Get signed download url.

Required Parameter:

- `gcs_key`: string

Possible Response Codes:

- 200 Successful Response
- 422 Validation Error

/api/v1/admin/users-with-uploads

Method:

GET

Description:

Fetch users with uploads.

Possible Response Codes:

- 200 Successful Response

Health Endpoints

/

Method:

GET

Description:

Health check endpoint. Used to verify that the Search Service is online and responsive.

Possible Response Codes:

- 200 Successful Response



VELOX

JUNE 2025

Search Service

The Search Service is responsible for querying the glossary database to retrieve relevant terms, suggest matching entries, and support health check functionality. This service powers both precise and fuzzy search across multiple filters including language and domain.

Search Endpoints

/api/v1/search/

Method:

GET

Description:

Searches database for terms matching given parameters.

Parameters:

- `query`: string
- `language`: string
- `domain`: string
- `sort_by`: name or popularity
- `page`: number
- `page_size`: number
- `fuzzy`: boolean

Possible Response Codes:

- 201 Successful Response
- 422 Validation Error

Suggest Endpoints



VELOX

JUNE 2025

/api/v1/suggest/

Method:

GET

Description:

Returns search suggestions based on a partial query using fuzzy matching.

Parameters:

- `query`: string

Possible Response Codes:

- 201 Successful Response
- 422 Validation Error

/api/v1/glossary/categories

Method:

GET

Description:

Retrieves all available glossary categories.

Possible Response Codes:

- 200 OK

/api/v1/glossary/categories/{category_name}/terms

Method:

GET

Description:

Fetches all terms associated with a specific category.

Path Parameters:

- `category_name`: string

Possible Response Codes:

- 200 OK



VELOX

JUNE 2025

- 422 Validation Error

/api/v1/glossary/terms/{term_id}/translations

Method:

GET

Description:

Fetches all translations for a specific term.

Path Parameters:

- `term_id`: string

Possible Response Codes:

- 200 OK
- 422 Validation Error

/api/v1/glossary/search

Performs a simple term search across all categories.

Query Parameters:

- `query`: string

Possible Response Codes:

- 200 OK
- 422 Validation Error

/api/v1/glossary/search

Method:

POST

Description:

Advanced search with filters (language, domain) and pagination.

Query Parameters:

- `query`: string



VELOX

JUNE 2025

- `domain`: string
- `language`: string
- `page`: integer (default: 1)
- `limit`: integer (default: 10)

Request Body:

```
{  
  "query": "string",  
  "domain": "string",  
  "language": "string"  
}
```

Possible Response Codes:

- 200 OK
- 422 Validation Error

/api/v1/glossary/domains

Method:

GET

Description:

Retrieves all available glossary domains.

Possible Response Codes:

- 200 OK

/api/v1/glossary/languages

Method:

GET

Description:

Retrieves all available glossary languages.

Possible Response Codes:



VELOX

JUNE 2025

- 200 OK

/api/v1/glossary/translate

Method:

POST

Description:

Translates a list of terms into one or more target languages.

Query Parameters:

- `source_language`: string (default: English)
- `domain`: string

Request Body:

```
{  
  "terms": ["string"],  
  "target_languages": ["string"]  
}
```

Possible Response Codes:

- 200 OK
- 422 Validation Error

/api/v1/glossary/stats

Method:

GET

Description:

Returns statistics about the glossary.

Possible Response Codes:

- 200 OK

/api/v1/glossary/random



VELOX

JUNE 2025

Method:

GET

Description:

Returns one or more random terms.

Query Parameters:

- `count`: integer (default: 1)

Possible Response Codes:

- 200 OK
- 422 Validation Error

Health Endpoints

/

Method:

Get

Description:

Health check endpoint. Used to verify that the Search Service is online and responsive.

Possible Response Codes:

- 201 Successful Response
- 500 Internal Server Error

Analytics Service

The Marito Analytics Service provides descriptive statistics and insights derived from the glossary data. It includes coverage metrics, length distributions, and summary reports for terms, languages, and categories.



VELOX

JUNE 2025

Analytics Endpoints

/api/v1/analytics/descriptive

Method:

GET

Description:

Returns all descriptive analytics in a combined legacy format.

Possible Response Codes:

- 200 OK

/api/v1/analytics/descriptive/category-frequency

Method:

GET

Description:

Returns frequency distribution of terms across different categories.

Possible Response Codes:

- 200 OK

/api/v1/analytics/descriptive/language-coverage

Method:

GET

Description:

Returns percentage of filled translations per language.

Possible Response Codes:

- 200 OK

/api/v1/analytics/descriptive/term-length

Method:

GET



VELOX

JUNE 2025

Description:

Returns average term lengths per language.

Possible Response Codes:

- 200 OK

/api/v1/analytics/descriptive/definition-length**Method:**

GET

Description:

Returns average definition lengths per language.

Possible Response Codes:

- 200 OK

/api/v1/analytics/descriptive/unique-terms**Method:**

GET

Description:

Returns the number of unique terms per language.

Possible Response Codes:

- 200 OK

/

Method:

GET

Description:

Health check endpoint.

Possible Response Codes:

- 200 OK



VELOX

JUNE 2025

Linguist Service

The Linguist ApplicationService enables authenticated users to submit documentation for linguist verification. This includes uploading identity documents, CVs, certifications, and optional research papers.

Linguist Endpoints

`/api/v1/linguist-applications/`

Method:

POST

Description:

Creates a new linguist application for the currently authenticated user.

Request Body:

```
json
CopyEdit
{
  "id_document_url": "string",
  "cv_document_url": "string",
  "certifications_document_url": "string",
  "research_papers_document_url": "string"
}
```

Possible Response Codes:

- 201 Created
- 422 Validation Error

Vote Service



VELOX

JUNE 2025

The Vote Service enables users to cast votes on glossary terms. Each term supports upvotes and downvotes per user, and vote counts are recalculated in real-time.

Vote Endpoints

/api/v1/votes/

Method:

POST

Description:

Casts or updates a vote for a specific term by the authenticated user. Sends the updated vote count in the response.

Request Body:

```
{  
  "term_id": "3fa85f64■5717■4562-b3fc-2c963f66afa6",  
  "vote": "upvote"  
}
```

Possible Response Codes:

- 200 OK

```
{  
  "term_id": "3fa85f64■5717■4562-b3fc-2c963f66afa6",  
  "upvotes": 0,  
  "downvotes": 0,  
  "user_vote": "upvote"  
}
```

- 422 Validation Error



VELOX

JUNE 2025